

Name: Renu Mariam Mathew

Student Id:16376

Hw#4

Step 1:

Create external table emp(eid int,name String,role String,salary String) row format delimited fields terminated by '\t' lines terminated by '\n';

Step 2:

```
LOAD DATA LOCAL INPATH 'cs570D.txt'  
OVERWRITE INTO TABLE EMP;
```

Step 1 and Step 2 combined

Name: Renu Mariam Mathew

Student Id:16376

Hw#4

```
hadoop@localhost:~  
OK  
Time taken: 1.13 seconds  
hive> Create external table emp(eid int,name String,role String,salary String)  
row format delimited fields terminated by '\t' lines terminated by '\n';  
OK  
Time taken: 1.292 seconds  
hive> LOAD DATA LOCAL INPATH 'cs570D.txt' OVERWRITE INTO TABLE EMP;  
Loading data to table default.emp  
Table default.emp stats: [numFiles=1, numRows=0, totalSize=123, rawDataSize=0]  
OK  
Time taken: 4.093 seconds  
hive> SELECT * FROM EMP;  
OK  
1201      Gopal      45000      manager  
1202      Manisha  45000      reader  
1203      Mast      40000      writer  
1204      Kiran      40000      Hr  
1205      Kranthi  30000      Op  
Time taken: 0.477 seconds, Fetched: 5 row(s)
```

Step 3:

set hive.exec.dynamic.partition= true;

set hive.exec.dynamic.partition.mode=nonstrict;

```
hadoop@localhost:~  
hive> set hive.exec.dynamic.partition= true;  
hive> set hive.exec.dynamic.partition.mode=nonstrict;
```

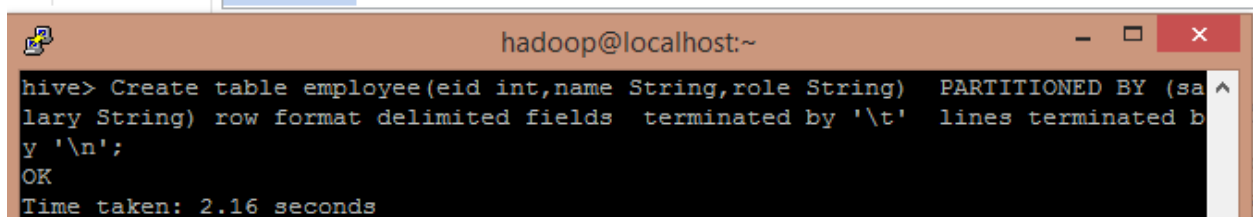
Name: Renu Mariam Mathew

Student Id:16376

Hw#4

Step 4:

Create table employee(eid int,name String,role String)
PARTITIONED BY (salary String) row format delimited
fields terminated by '\t' lines terminated by '\n';

A terminal window titled 'hadoop@localhost:~' with standard window controls. The terminal shows a Hive command being executed: 'hive> Create table employee(eid int,name String,role String) PARTITIONED BY (salary String) row format delimited fields terminated by '\t' lines terminated by '\n';'. The output shows 'OK' and 'Time taken: 2.16 seconds'.

```
hadoop@localhost:~  
hive> Create table employee(eid int,name String,role String) PARTITIONED BY (sa  
lary String) row format delimited fields terminated by '\t' lines terminated b  
y '\n';  
OK  
Time taken: 2.16 seconds
```

Step 5:

INSERT OVERWRITE TABLE EMPLOYEE
PARTITION (salary) select eid,name,salary,role from
emp;

Name: Renu Mariam Mathew

Student Id:16376

Hw#4

```
hive> set hive.exec.dynamic.partition.mode=nonstrict;
hive> INSERT OVERWRITE TABLE EMPLOYEE PARTITION (salary) select eid,name,salary,
role from emp;
Query ID = hadoop_20160405153737_22624c46-af17-4db0-b281-03c4d0818c1b
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1459881998495_0006, Tracking URL = http://localhost:8088/proxy/application_1459881998495_0006/
Kill Command = /home/hadoop/hadoop/bin/hadoop job -kill job_1459881998495_0006
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2016-04-05 15:38:25,360 Stage-1 map = 0%, reduce = 0%
2016-04-05 15:39:26,172 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 5.87 sec
MapReduce Total cumulative CPU time: 5 seconds 870 msec
Ended Job = job_1459881998495_0006
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
```

OutPut:

hdfs dfs -cat
/user/hive/warehouse/employee/salary=30000/000000_0

127.0.0.1:50070/explorer.html#/user/hive/warehouse/employee

Permission	Owner	Group	Size	Replication	Block Size	Name
drwxrwxr-x	hadoop	supergroup	0 B	0	0 B	salary=30000
drwxrwxr-x	hadoop	supergroup	0 B	0	0 B	salary=40000
drwxrwxr-x	hadoop	supergroup	0 B	0	0 B	salary=45000

/user/hive/warehouse/employee/salary=30000						
Permission	Owner	Group	Size	Replication	Block Size	Name
-rwxrwxr-x	hadoop	supergroup	16 B	1	128 MB	000000_0

Hadoop, 2014.

SnapshotStartup ProgressUtilities

File information - 000000_0

Download

Block information --

Block ID: 1073742548
Block Pool ID: BP-189
Generation Stamp: 1
Size: 16
Availability:

- localhost

000000_0

File Edit Format View Help

1205Kranthi Op

Name: Renu Mariam Mathew

Student Id:16376

Hw#4

The input file which I used is

```
1201 Gopal 45000 manager
1202 Manisha 45000 reader
1203 Mast 40000 writer
1204 Kiran 40000 Hr
1205 Kranthi 30000 Op
```